

Algorithmic Thinking

Luay Nakhleh

Clustering and the Closest Pair Problem

The Divide-and-Conquer Algorithmic Technique

1 Clustering

Definition 1 A clustering of a set P of points into k clusters is a partition of P into sets C_1, \dots, C_k , such that

- $\forall 1 \leq i \leq k, C_i \subseteq P$,
- $\forall 1 \leq i \leq k, C_i \neq \emptyset$,
- $\forall 1 \leq i, j \leq k, i \neq j, C_i \cap C_j = \emptyset$, and
- $\cup_{i=1}^k C_i = P$.

Fig. 1 shows a clustering of 10 points into two clusters.

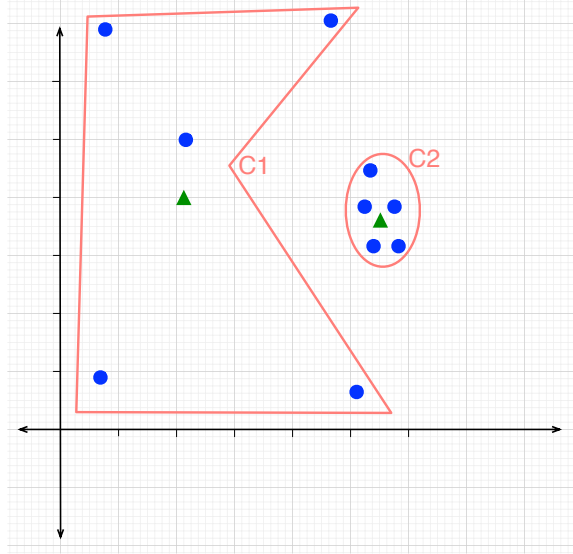


Figure 1: Two clusters C_1 and C_2 on the set of points (blue solid circles) and their centers are shown (green triangle).

We define the center of a cluster C_u as

$$\text{center}(C_u) = \frac{1}{|C_u|} \sum_{p_i \in C_u} (x_i, y_i).$$

For example, if $C_u = \{p_1, p_7, p_9\}$, with $p_1 = (1, 2)$, $p_2 = (4, 6)$, and $p_3 = (4, 4)$, then

$$\text{center}(C_u) = \frac{1}{3}((1, 2) + (4, 6) + (4, 4)) = \frac{1}{3}(9, 12) = (3, 4).$$

Fig. 1 shows the centers of the two clusters C1 and C2.

While many clusterings of the points in P exist, a desired property is that the partition results in clusters with higher similarity of points within a cluster than of points between clusters. Algorithms **HierarchicalClustering** and **KMeansClustering** below are two heuristics for generating clustering with this desired property. In both algorithms, we will use k to denote the number of clusters.

A word on implementation. While P is defined as a set in both clustering algorithms, it is more convenient to implement it using a list, since each point can be accessed directly in the list. Further, both algorithms return a set C of clusters; here, C is a set of elements, where each element is a set of points, and C satisfies the properties in Definition 1.

Algorithm 1: HierarchicalClustering.

Input: A set P of points whose i th point, p_i , is a pair (x_i, y_i) ; k , the desired number of clusters.

Output: A set C of k clusters that provides a clustering of the points in P .

```

1  $n \leftarrow |P|$ ;
2 Initialize  $n$  clusters  $C = \{C_1, \dots, C_n\}$  such that  $C_i = \{p_i\}$ ;
3 while  $|C| > k$  do
4    $(C_i, C_j) \leftarrow \operatorname{argmin}_{C_i, C_j \in C, i \neq j} d_{C_i, C_j}$ ;
5    $C \leftarrow C \cup \{C_i \cup C_j\}$ ;
6    $C \leftarrow C \setminus \{C_i, C_j\}$ ;
7 return  $C$ ;
```

Algorithm 2: KMeansClustering.

Input: A set P of points whose i th point, p_i , is a pair (x_i, y_i) ; k , the desired number of clusters; q , a number of iterations.

Output: A set C of k clusters that provides a clustering of the points in P .

```

1  $n \leftarrow |P|$ ;
2 Initialize  $k$  centers  $\mu_1, \dots, \mu_k$  to initial values (each  $\mu$  is a point in the 2D space);
3 for  $i \leftarrow 1$  to  $q$  do
4   Initialize  $k$  empty sets  $C_1, \dots, C_k$ ;
5   for  $j = 0$  to  $n - 1$  do
6      $\ell \leftarrow \operatorname{argmin}_{1 \leq f \leq k} d_{p_j, \mu_f}$ ;
7      $C_\ell \leftarrow C_\ell \cup \{p_j\}$ ;
8   for  $f = 1$  to  $k$  do
9      $\mu_f = \operatorname{center}(C_f)$ ;
10 return  $\{C_1, C_2, \dots, C_k\}$ ;
```

2 2D Points, the Euclidian Distance, and Cluster Error

Both algorithms, **HierarchicalClustering** and **KMeansClustering**, make use of a distance measure, d . In the case of **HierarchicalClustering**, d_{C_i, C_j} is the distance between the two clusters C_i and C_j . In the case of **KMeansClustering**, d_{p_j, μ_f} is the distance between the point p_j and center μ_f of cluster C_f . But, how is this distance measure d defined?

In our case, we will only deal with points in the 2D space, such that each point p_i is given by two features: its horizontal (or, x) and vertical (or, y) coordinates, so that $p_i = (x_i, y_i)$. One natural way to quantify the distance between two points p_i and p_j in this case is the standard Euclidian distance:

$$d_{p_i, p_j} = \sqrt{(x_i - x_j)^2 + (y_i - y_j)^2}.$$

While this distance measure applies directly to compute d_{p_j, μ_f} in **KMeansClustering**, how does it apply to computing d_{C_i, C_j} in **HierarchicalClustering**? By computing the distance between their centers:

$$d_{C_i, C_j} \equiv d_{center(C_i), center(C_j)}.$$

In our discussion above, we assumed a known number of clusters, k , and sought a clustering with that number of clusters. However, in general, the number of clusters is unknown. One way to determine the number of clusters is to vary the value of k , such that $k = 1, 2, 3, \dots$, and for each value of k to “inspect” the quality of the clusters obtained. One measure of such quality is the error of a cluster, which reflects how tightly packed around the center the cluster’s points are, and is defined for cluster C_i as

$$error(C_i) = \sum_{p \in C_i} (d_{p, center(C_i)})^2.$$

To illustrate, consider the two clusters in Fig. 1. Cluster C1 has a larger *error* value than C2 and, indeed, compared to C2, it is hard to argue that the points in C1 form a cluster.

3 The Closest Pair Problem

Now that we have defined the distance between points and clusters, we need an algorithm that finds, among a set of clusters, two clusters that are closest to each other (in the case of **HierarchicalClustering**), or, among a set of centers, a closest center to a given point (in the case of **KMeansClustering**). In this Module, we will approach this task by solving the Closest Pair problem, defined as follows:

- **Input:** A set P of (distinct) points and a distance measure d defined on every two points in P .
- **Output:** A pair of distinct points in P that are closest to each other under the distance measure d .

Notice that the solution of the problem might not be unique (that is, more than a single pair of points might be closest), in which case we are interested in an arbitrary one of those pairs with the smallest pairwise distance.

A simple brute-force algorithm can solve this problem, as given by the pseudo-code of Algorithm **SlowClosestPair**. Notice the notation we use for finding the minimum of two tuples $\min\{(d_1, p_1, q_1), (d_2, p_2, q_2)\}$, which returns the tuple that has the smallest first element (that is, it returns tuple (d_1, p_1, q_1) if $d_1 < d_2$, and (d_2, p_2, q_2) otherwise). In case the two tuples have the same first element, one of them is returned arbitrarily.

Algorithm 3: SlowClosestPair.

Input: A set P of (≥ 2) points whose i th point, p_i , is a pair (x_i, y_i) .

Output: A tuple (d, i, j) where d is the smallest pairwise distance of points in P , and i, j are the indices of two points whose distance is d .

```

1  $(d, i, j) \leftarrow (\infty, -1, -1);$ 
2 foreach  $p_u \in P$  do
3   foreach  $p_v \in P$  ( $u \neq v$ ) do
4      $(d, i, j) \leftarrow \min\{(d, i, j), (d_{p_u, p_v}, u, v)\};$  // min compares the first element of each tuple
5 return  $(d, i, j);$ 
```

Can we do better than **SlowClosestPair** in terms of running time? We will now consider a divide-and-conquer algorithm for this problem, **FastClosestPair**.

Algorithm 4: FastClosestPair.

Input: A set P of (≥ 2) points whose i th point, p_i , is a pair (x_i, y_i) , **sorted** in nondecreasing order of their horizontal (x) coordinates.

Output: A tuple (d, i, j) where d is the smallest pairwise distance of the points in P , and i, j are the indices of two points whose distance is d .

```

1  $n \leftarrow |P|$ ;
2 if  $n \leq 3$  then
3    $(d, i, j) \leftarrow \text{SlowClosestPair}(P)$ ;
4 else
5    $m \leftarrow \lfloor n/2 \rfloor$ ;
6    $P_\ell \leftarrow \{p_i : 0 \leq i \leq m-1\}$ ;  $P_r \leftarrow \{p_i : m \leq i \leq n-1\}$ ;           //  $P_\ell$  and  $P_r$  are also sorted
7    $(d_\ell, i_\ell, j_\ell) \leftarrow \text{FastClosestPair}(P_\ell)$ ;
8    $(d_r, i_r, j_r) \leftarrow \text{FastClosestPair}(P_r)$ ;
9    $(d, i, j) \leftarrow \min\{(d_\ell, i_\ell, j_\ell), (d_r, i_r + m, j_r + m)\}$ ;
10   $mid \leftarrow \frac{1}{2}(x_{m-1} + x_m)$ ;                                           // center line of strip
11   $(d, i, j) \leftarrow \min\{(d, i, j), \text{ClosestPairStrip}(P, mid, d)\}$ ;
12 return  $(d, i, j)$ ;
```

Algorithm 5: ClosestPairStrip.

Input: A set P of points whose i th point, p_i , is a pair (x_i, y_i) ; mid and w , both of which are real numbers.

Output: A tuple (d, i, j) where d is the smallest pairwise distance of points in P whose horizontal (x) coordinates are within w from mid .

```

1 Let  $S$  be a list of the set  $\{i : |x_i - mid| < w\}$ ;
2 Sort the indices in  $S$  in nondecreasing order of the vertical ( $y$ ) coordinates of their associated points;
3  $k \leftarrow |S|$ ;
4  $(d, i, j) \leftarrow (\infty, -1, -1)$ ;
5 for  $u \leftarrow 0$  to  $k - 2$  do
6   for  $v \leftarrow u + 1$  to  $\min\{u + 3, k - 1\}$  do
7      $(d, i, j) \leftarrow \min\{(d, i, j), (d_{p_{S[u]}, p_{S[v]}}, S[u], S[v])\}$ ;
8 return  $(d, i, j)$ ;
```